



EARL
CONFERENCE

LONDON 2018

Scaling R in the Cloud

with Azure Batch and Docker

Christoph Bodner



01

02

03

Topics

Who we are

(obligatory marketing stuff...)

Data Science@Post AG:

- Overview: Post AG
- Our team

Our problems

Deploying R at scale:

- Scaling compute resources on demand
- Reduce dev-prod disparity

Our solution

Main components:

- Azure Batch
- Docker
- VSTS



01

02

03

Topics

Who we are

(obligatory marketing stuff...)

Data Science@Post AG:

- Overview: Post AG
- Our team

Our problems

Deploying R at scale:

- Scaling compute resources on demand
- Reduce dev-prod disparity

Our solution

Main components:

- Azure Batch
- Docker
- VSTS




OVERVIEW: POST AG

COMPANY PERFORMANCE & PARCEL VOLUMES OVER TIME


SALES PERFORMANCE

mio EUR



PARCEL VOLUMES OF AUSTRIAN POST

mio parcels



OVERVIEW: POST AG

STRONG PRESENCE IN EASTERN EUROPE

GROWTH FOCUS ON PARCEL & LOGISTICS DIVISION



ORGANISATION

DATA SCIENCE IS PART OF BI COMPETENCE CENTER



Lead BI CC  **Post**
Georg Posan

BI Competence Center



SAP Gurus

Everything related to
SAP HANA, BO, etc.



Microsoft

.net/T-SQL Wizards

SQL Server, PowerBI and
MS tech in general



This is us ☺

Analytics team



OUR TEAM

PEOPLE WHO LIKE $\pi z^2 a$ IN EVERY FORM



Christoph Bodner

Lead Data Scientist

Quantitative Finance (WU)
Prev.: KPMG



Thomas Laber

Senior Data Scientist

Business Informatics (TU)
Prev.: Accenture



Martin Blöschl

Junior Data Scientist

Computational Intelligence (TU)



Raphael Pesl

Junior Data Scientist

Mathematics (TU)



01

02

03

Topics

Who we are

(obligatory marketing stuff...)

Data Science@Post AG:

- Overview: Post AG
- Our team

Our problems

Deploying R at scale:

- Scaling compute resources on demand
- Reduce dev-prod disparity

Our solution

Main components:

- Azure Batch
- Docker
- VSTS

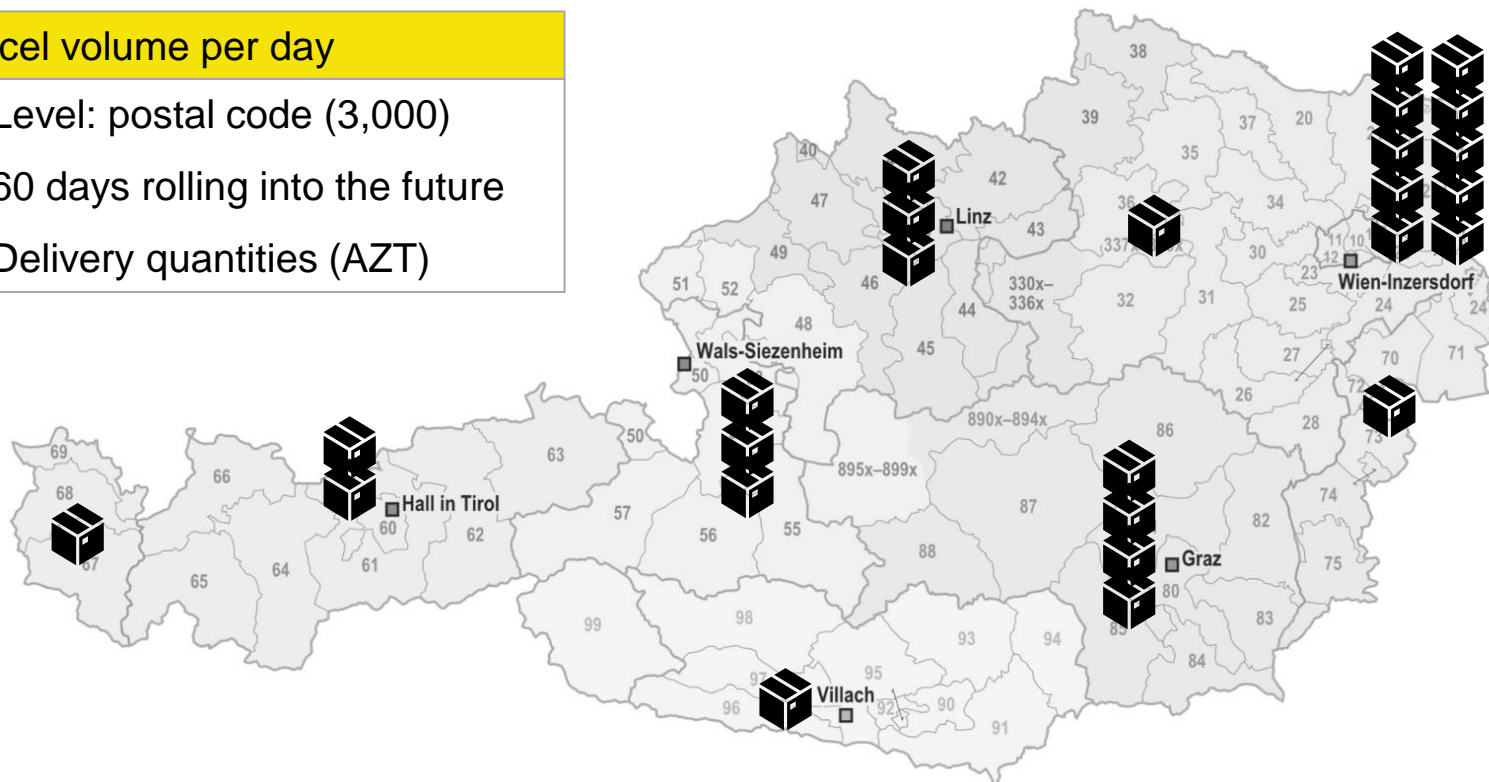


PARCEL VOLUME FORECAST


HOW MANY PARCELS WILL WE NEED TO DELIVER IN THE FUTURE?

Parcel volume per day

- Level: postal code (3,000)
- 60 days rolling into the future
- Delivery quantities (AZT)

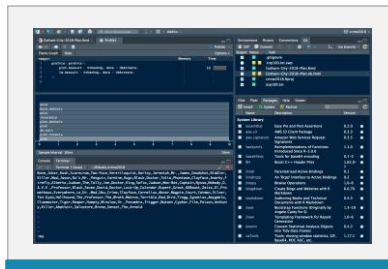


FROM EXPERIMENT TO DEPLOYMENT OUR JOURNEY



Whiteboard phase

- Draft architecture
- Modeling strategy
- Everything seems easy 😊



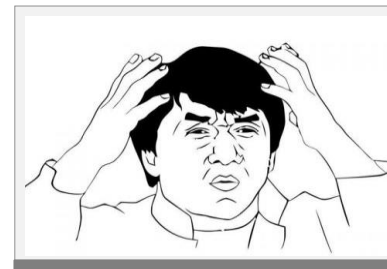
Experiment phase

- Build model on sample
- Model improves
- Happy Data Scientists



Deadline phase

- Model training takes too long
- Model gets simplified to speed up training
- Performance suffers



Deployment phase

- Dev \neq Prod
- Server procurement slow
- Server not powerful enough for training
- Server too powerful for scoring




FORECASTING PERFORMANCE

2-DAYS-AHEAD


Performance

Parcel volumes per day Austria (2 days ahead)

Simplified model



Deviation in percent

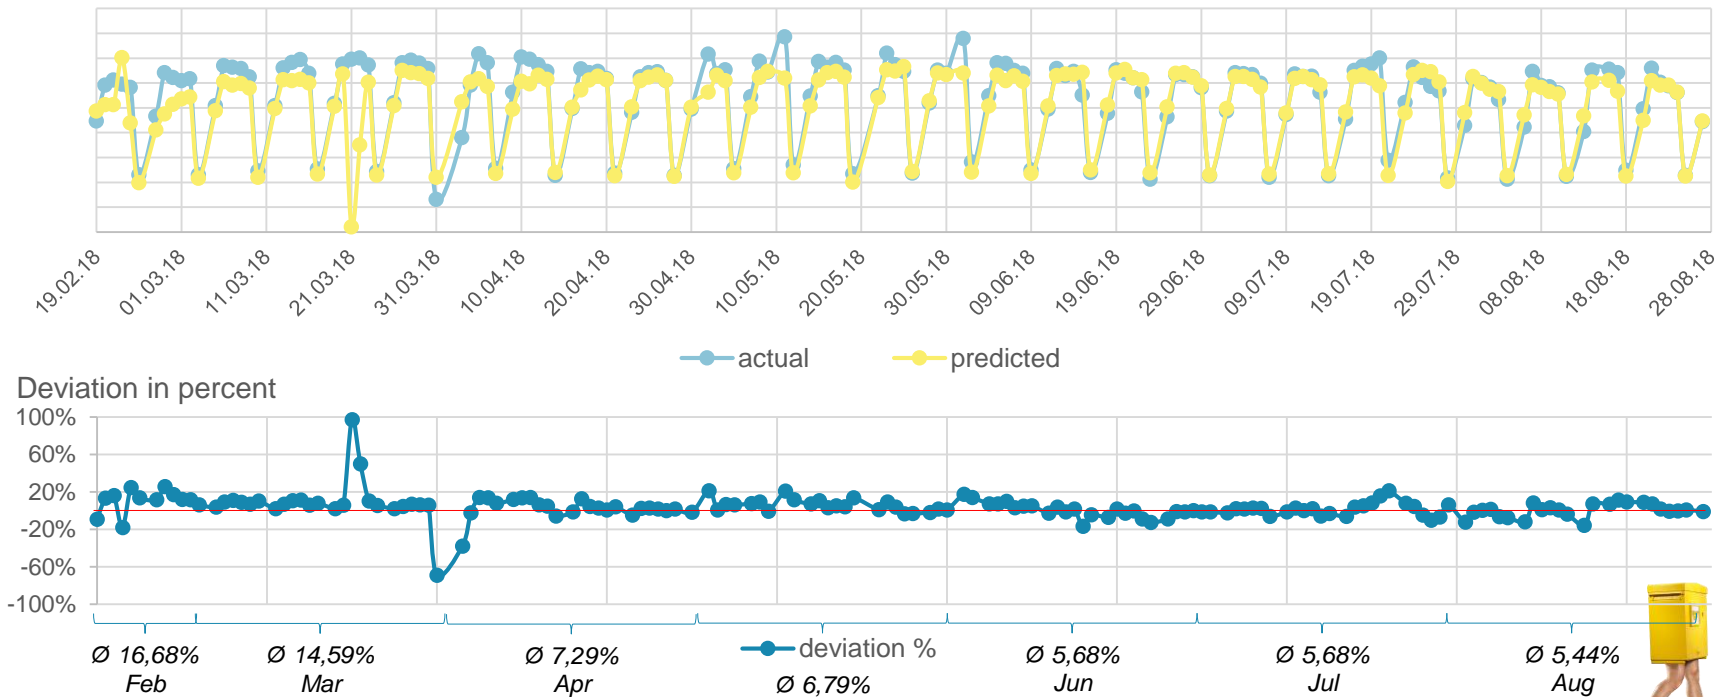


FORECASTING PERFORMANCE

7-DAYS-AHEAD

Performance


Parcel volumes per day Austria (7 days ahead)




PROBLEMS IN DETAIL

SCALING COMPUTE RESOURCES ON DEMAND

Resource utilisation



Scaling options




► We want to be able to elastically scale **up** and **out** to meet our needs



PROBLEMS IN DETAIL

DEV-PROD DISPARITY

Happy Path Deployments



Happy Path Deployments



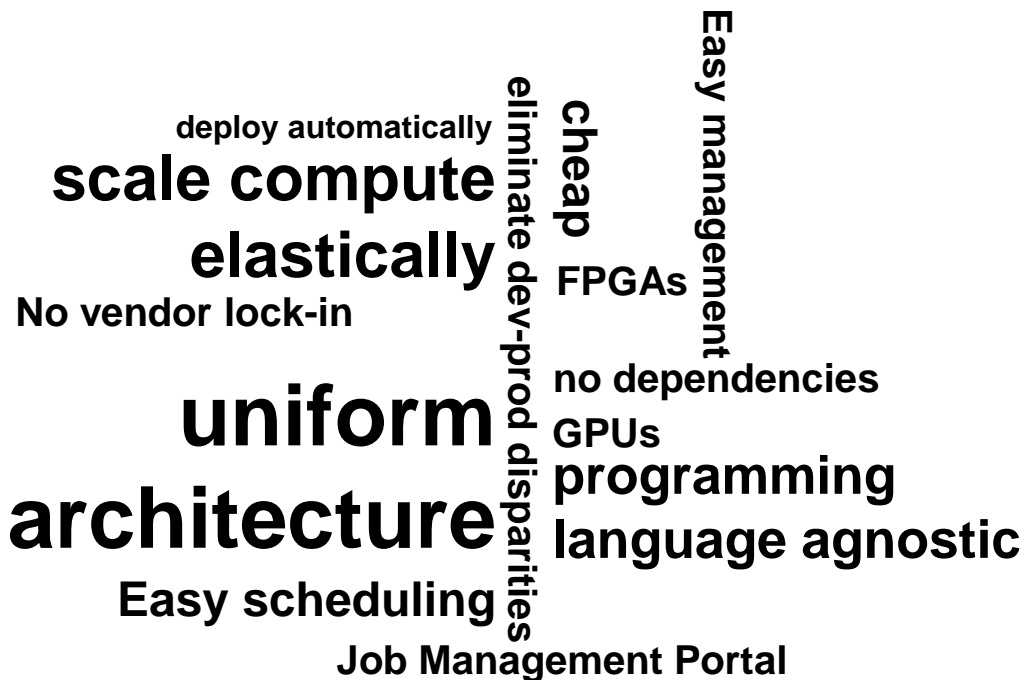
► Keeping Dev and Prod aligned is hard even with best intentions



HOW CAN WE SOLVE OUR PROBLEMS?

WE HAVE A LONG WISH LIST

Our wish list:



A word cloud of requirements for solving problems. The words are arranged in a circular pattern, with 'uniform architecture' and 'scale compute elastically' being the largest. Other words include 'cheap', 'FPGAs', 'no dependencies', 'programming language agnostic', 'Easy scheduling', 'Job Management Portal', 'Easy management', 'eliminate dev-prod disparities', 'No vendor lock-in', and 'deploy automatically'.

uniform architecture

scale compute elastically

cheap

FPGAs

no dependencies

programming language agnostic

Easy scheduling

Job Management Portal

Easy management

eliminate dev-prod disparities

No vendor lock-in

deploy automatically

Solutions we considered:

We have two dominant stacks:

- Microsoft .net + Azure
- SAP

Since we have SAP and .net/Azure support in-house, we looked primarily at these two stacks.



01

02

03

Topics

Who we are

(obligatory marketing stuff...)

Data Science@Post AG:

- Overview: Post AG
- Our team

Our problems

Deploying R at scale:

- Scaling compute resources on demand
- Reduce dev-prod disparity

Our solution

Main components:





- Docker
- Azure Batch
- VSTS



SOLUTIONS WE CONSIDERED COMPARISON

Also take a look at
Azure Databricks for
Spark workloads

Cloud vs. On-Premise Solution

	Azure ML Studio	Azure ML Workbench	Azure Batch	SAP HANA (PAL)
Supported Languages	R/Python	Python	All	R (Python)
Supports GPUs	This looks very promising	Yes	Yes	No
Dev-Prod parity	Easy	Easy	Easy	Hardish
R package available	AzureML	-	doAzureParallel	-
Independent upgrades	partially	partially	Yes	partially
Elastic scaling	-	Yes	Yes	No
Scheduling included	No	No	Yes	Yes
SCORE				




AZURE BATCH OVERALL WINNER

FLEXIBLE, HIGH VALUE/MONEY AND LOW VENDOR LOCK-IN




Scale up and out

Specify node sizes and types, e.g. GPU/CPU, RAM and get large discount on low-prio nodes



Scheduling integrated

Specify job schedule and resize pool based on number of outstanding tasks



Docker support

Dev and Prod parity
Fits into CD pipeline



DOCKER BRIEFLY EXPLAINED

BUILD ONCE – RUN ANYWHERE


Example Dockerfile

```
1  ## Description: https://hub.docker.com/r/rocker/tidyverse/
2  FROM rocker/tidyverse
3
4  # Install your R package
5  ## Copy R package to docker
6  RUN mkdir -p /usr/r_package
7  COPY . /usr/r_package
8
9  ## Install package dependencies
10 RUN R CMD SHLIB /usr/r_package
11 ## Roxygenize
12 RUN R CMD SHLIB /usr/r_package
13 ## Install package
14 RUN R CMD INSTALL /usr/r_package
```

R image used

Here we install our package

Build Container

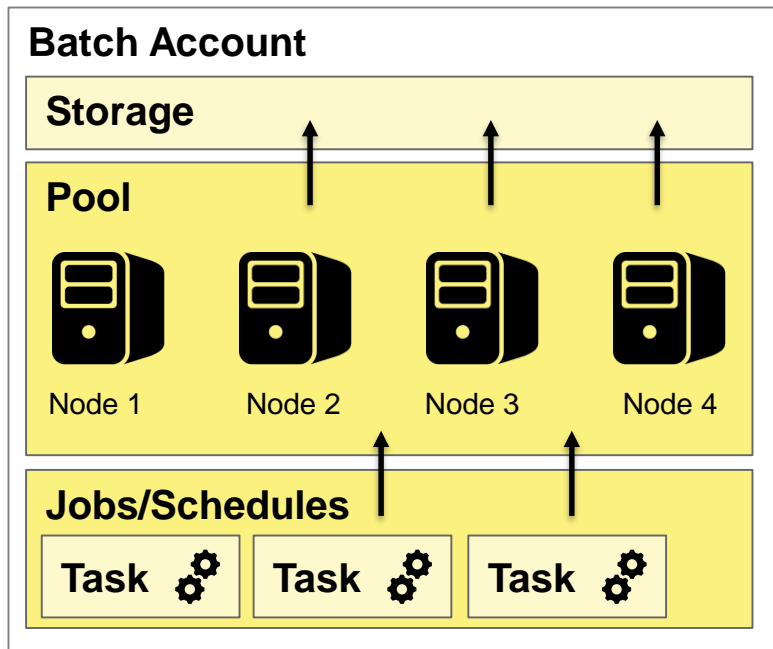


► Container can be deployed to docker runtime in production: Dev = Prod (mostly☺)

Source image: 1) <https://stephenafamo.com/blog/understanding-docker-containers/>



Components overview



Description

Pool Definition of compute resources

- Node number/type
- Container settings
- Autoscaling
- Task number per node

Schedule ... Recurring job for a pool

- Job specification

Job Contains pool and task specs

- Job manager task

Task Work specification

- Command line
- upload options



AZURE BATCH SETUP OVERVIEW

Pool configuration

```
1 # setup batch pool
2 def create_pool(batch_service_client, pool_id)
3     image_ref_to_use = batch.models.ImageReference(
4         publisher='microsoft-azure-batch',
5         offer='ubuntu-server-container',
6         sku='16-04-lts',
7         version='latest'
8     )
9
10    # Private docker repo specification
11    container_conf = batch.models.ContainerConfiguration(
12        ...
13    )
14
15    # create pool specs with pre-fetched images
16    new_pool = batch.models.PoolAddParameter(
17        id=pool_id,
18        virtual_machine_configuration=batch.models.VirtualMachineConfiguration(
19            image_reference=image_ref_to_use,
20            container_configuration=container_conf,
21            node_agent_sku_id='batch.node.ubuntu 16.04',
22            vm_size=_POOL_VM_SIZE,
23            max_tasks_per_node=_MAX_TASKS_PER_NODE
24            enable_auto_scale=True,
25            auto_scale_evaluation_interval='PT5M',
26            auto_scale_formula= "some formula for node type and count"
27        )
28
29    # submit pool creation to azure batch
30    batch_service_client.pool.add(new_pool)
31
```

Pool base image

Here using
Python SDK

Docker image
& autoscaling

Job Manager Task

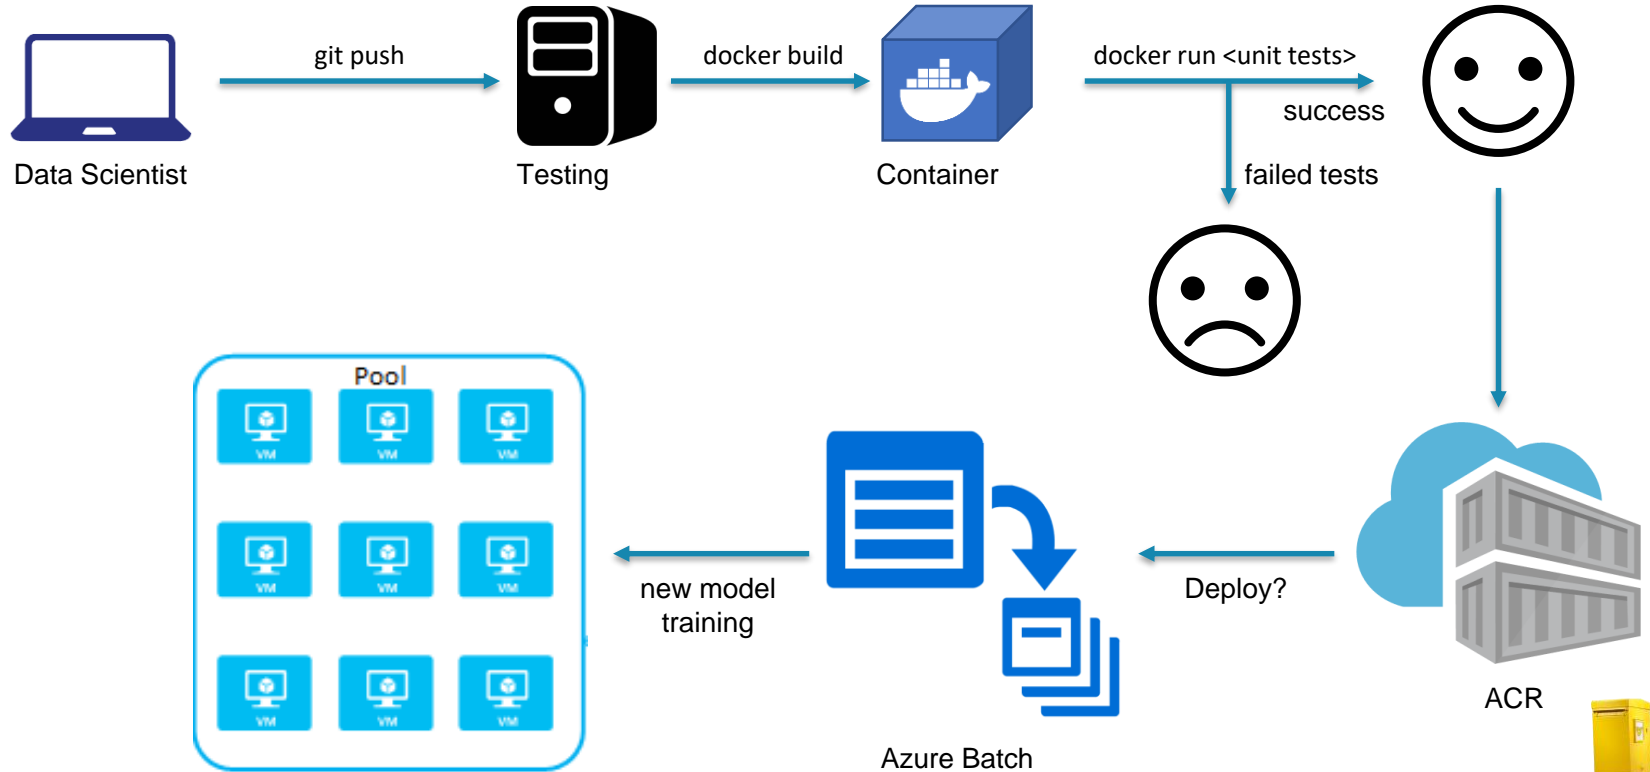
```
33 job_manager = batchmodels.JobManagerTask(
34     id=_AZ_JOB_MANAGER_ID,
35     command_line=_JOB_MANAGER_CMD,
36     container_settings=batch.models.ContainerSettings(
37         authentication_token_settings=batch.models.AuthenticationTokenSettings(
38             access=[batchmodels.AccessScope(
39                 ...
40             )],
41             environment_settings=[
42                 ...
43             ]
44         ),
45         kill_job_on_completion=True,
46         run_exclusive=False
47     )
48
49 job_spec = batchmodels.JobSpecification(
50     pool_info=batchmodels.PoolInformation(pool_id=_POOL_ID),
51     display_name='test_schedule',
52     #on_all_tasks_completion='noaction',
53     job_manager_task=job_manager,
54     #common_environment_settings=None
55 )
56
57 schedule = batchmodels.Schedule(
58     recurrence_interval='PT1M'
59 )
60
61 job_schedule = batchmodels.JobScheduleAddParameter(
62     id=_JOB_SCHEDULE_ID,
63     schedule=schedule,
64     job_specification=job_spec
65 )
66
67 # submit job schedule creation
68 batch_client.job_schedule.add(job_schedule)
```

Command to run

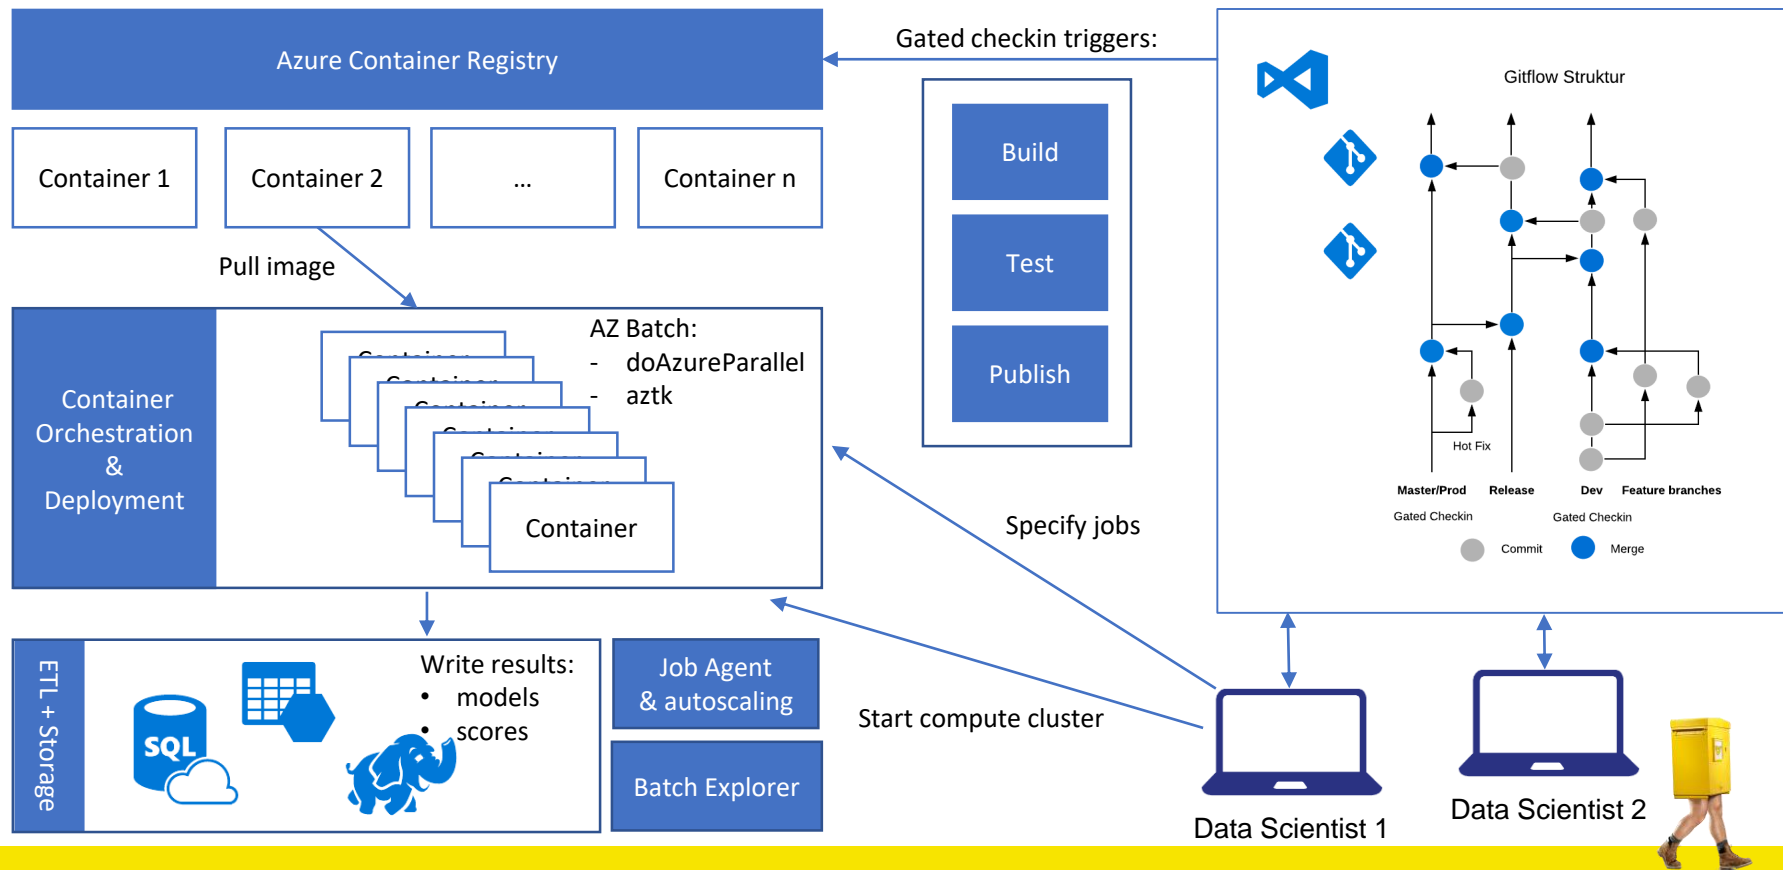
Job schedule
bundles everything



CONTINUOUS DELIVERY PIPELINE (SIMPLIFIED)

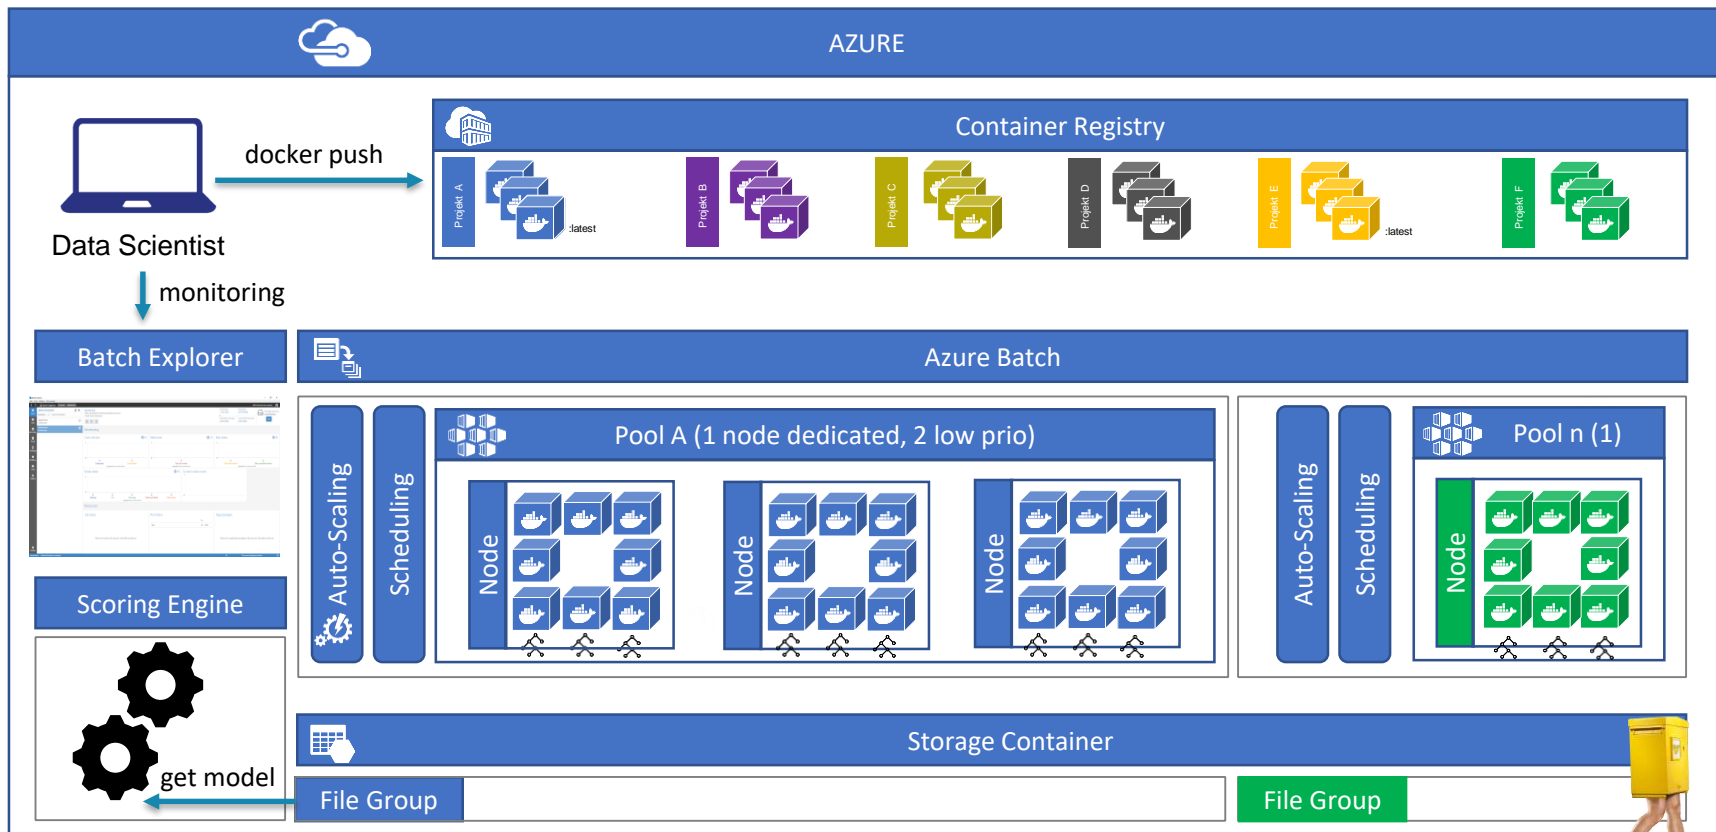


OUR STACK OVERVIEW

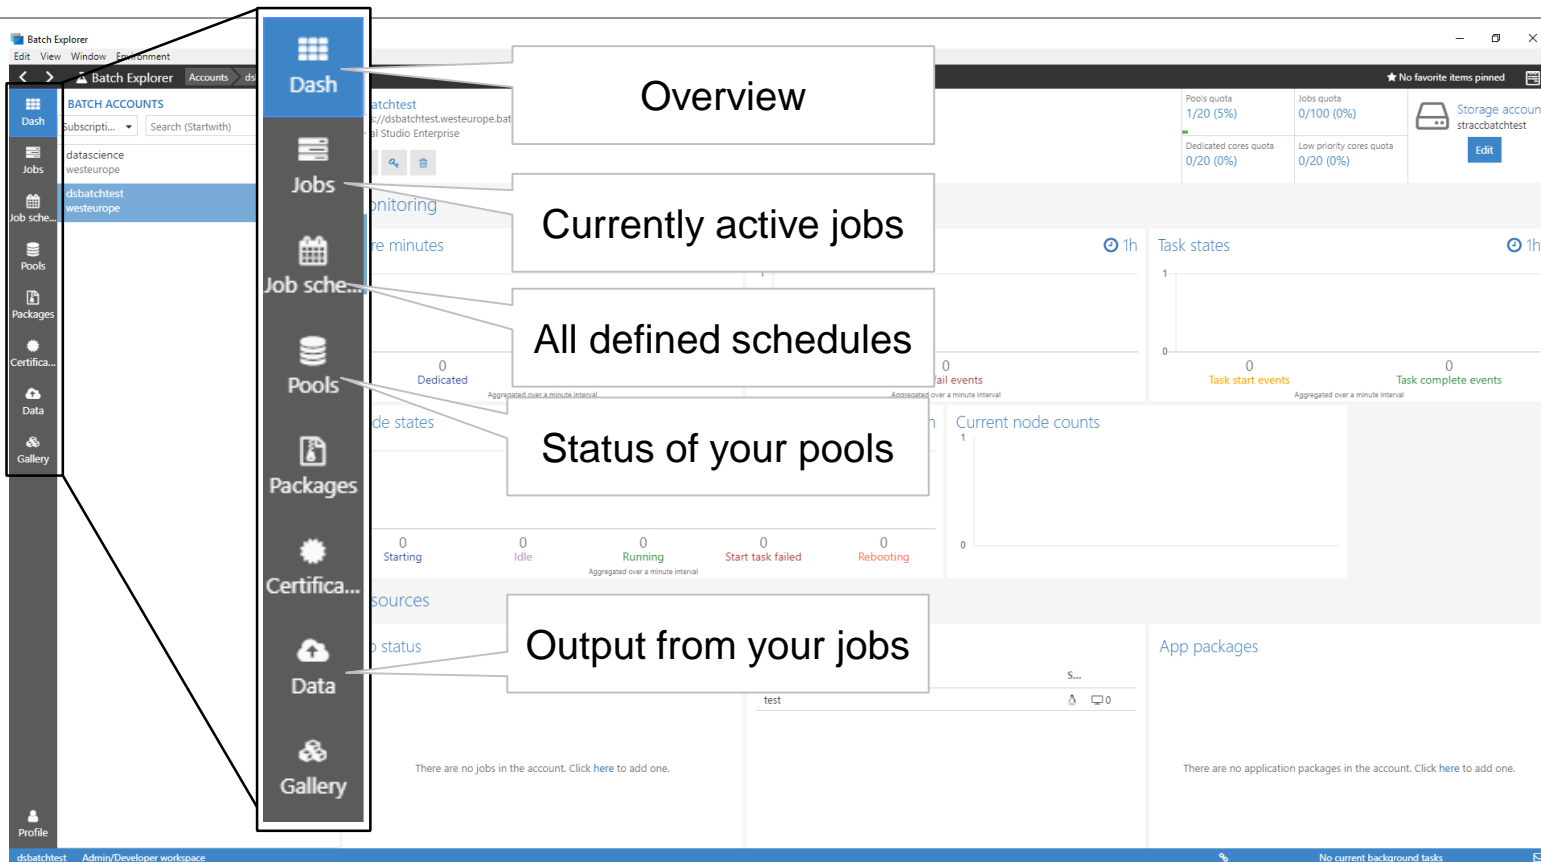


OUR STACK

WE ARE CURRENTLY BUILDING A HPC ENVIRONMENT



MONITORING BATCH JOBS WITH BATCH EXPLORER



The screenshot shows the Batch Explorer web application interface. A vertical sidebar on the left contains icons for various sections: Dash, Jobs, Job schedules, Pools, Packages, Certifications, Data, and Gallery. The main content area displays the 'Overview' section, which includes a top navigation bar, a 'BATCH ACCOUNTS' section, and several monitoring charts and tables. Callouts point to specific sections of the interface:

- Overview**: Points to the top navigation bar.
- Currently active jobs**: Points to the 'Jobs' section in the sidebar.
- All defined schedules**: Points to the 'Job schedules' section in the sidebar.
- Status of your pools**: Points to the 'Pools' section in the sidebar.
- Output from your jobs**: Points to the 'Data' section in the sidebar.

The main content area displays the 'Overview' section, which includes a top navigation bar, a 'BATCH ACCOUNTS' section, and several monitoring charts and tables. The 'BATCH ACCOUNTS' section shows a list of accounts, including 'dsbatchtest.westeurope.batch.azure.com'. The 'Overview' section includes a top navigation bar, a 'BATCH ACCOUNTS' section, and several monitoring charts and tables. The 'Pools' section shows a table of pools with columns for 'Pools quota', 'Jobs quota', 'Dedicated cores quota', and 'Low priority cores quota'. The 'Task states' section shows a line chart of task states over time. The 'Current node counts' section shows a line chart of node counts over time. The 'App packages' section shows a table of application packages.



SECURITY & EVALUATING DIFFERENT REAL-TIME SCORING OPTIONS

Security

- ~~Secret management with KeyVault~~
- ~~Azure Active Directory integration~~
- ~~VNET integration of pools~~
- Docker security best practices audit/check
- Disable public endpoints

What we are currently looking into

Real-time scoring options

- Using AzureML Studio
- Using Azure Model Management
- Kubernetes Cluster
- Azure Functions
- Azure Container Instances + Logic Apps
- ...



A bit of work is still open, but we plan to have everything wrapped up end of September



Thank you for
your attention!

Any questions?



Feel free to reach out to me:



christoph.bodner@post.at



[linkedin.com/in/christoph-bodner](https://www.linkedin.com/in/christoph-bodner)

